RISC-V Processor Trace
Version 0.029-DRAFT
f389d4381ac8edac1c2a076098df18f8dd293cb9

Gajinder Panesar, Iain Robertson

<gajinder.panesar@ultrasoc.com>, <iain.robertson@ultrasoc.com>

UltraSoC Technologies Ltd.

October 7, 2019

Contents

List of Figures

List of Tables

Chapter 1

Introduction

In complex systems understanding program behavior is not easy. Unsurprisingly in such systems, software sometimes does not behave as expected. This may be due to a number of factors, for example, interactions with other cores, software, peripherals, realtime events, poor implementations or some combination of all of the above.

It is not always possible to use a debugger to observe behavior of a running system as this is intrusive. Providing visibility of program execution is important. This needs to be done without swamping the system with vast amounts of data.

One method of achieving this is via a Processor Branch Trace.

This works by tracking execution from a known start address and sending messages about the address deltas taken by the program. These deltas are typically introduced by jump, call, return and branch type instructions, although interrupts and exceptions are also types of deltas.

Conceptually, the system has one or more of the following fundamental components:

- A core with an instruction trace interface that outputs all relevant information to successfully create a processor branch trace and more. This is a high bandwidth interface: in most implementations, it will supply a large amount of data (instruction address, instruction type, context information, ...) for each core execution clock cycle.

- A hardware encoder that takes in the CPU instruction trace and compresses it into lower bandwidth trace packets.

- A transmission channel to transmit or a memory to store these trace packets.

- A decoder, usually software on an external PC, that takes in the trace packets and, with knowledge of the program binary that's running on the originating hart, reconstructs the program flow. This decoding step can be done off-line or in real-time while the hart is executing.

In RISC-V, all instructions are executed unconditionally or at least their execution can be determined based on the program binary. The instructions between the deltas can all be assumed to

be executed sequentially. Because of this, there is no need to report sequential instructions in the trace, only whether the branches were taken or not and the address of taken indirect branches or jumps. If the program counter is changed by an amount that cannot be determined from the execution binary, the trace decoder needs to be given the destination address (i.e. the address of the next valid instruction). Examples of this are indirect branches or jumps, where the next instruction address is determined by the contents of a register rather than a constant embedded in the source code.

Interrupts generally occur asynchronously to the program's execution rather than intentionally as a result of a specific instruction or event. Exceptions can be thought of in the same way, even though they can be typically linked back to a specific instruction address. The decoder generally does not know where an interrupt occurs in the instruction sequence, so the trace encoder must report the address where normal program flow ceased, as well as give an indication of the asynchronous destination which may be as simple as reporting the exception type. When an interrupt or exception occurs, or the processor is halted, the final instruction executed beforehand must be included in the trace.

This document serves to specify the ingress port (the signals between the RISC-V core and the encoder), compressed branch trace algorithm and the packet format used to encapsulate the compressed branch trace information.

1.0.1   Nomenclature

In the following sections items in **bold** are signals or attributes within a packet.

Items in *italics* refer to parameters either built into the hardware or configurable hardware values.

An encoder is a piece of hardware that takes in RISC-V instruction data on its ingress port and transforms it into trace packets.

A decoder is a piece of software that takes the trace packets emitted by the encoder and reconstructs the execution flow of the code executed in the RISC-V core.

RISC-V has the following definitions:

- **Exception**: an unusual condition occurring at run time associated with an instruction in the current RISC-V hart

- **Interrupt**: an external asynchronous event that may cause a RISC-V hart to experience an unexpected transfer of control

- **Trap**: the transfer of control to a trap handler caused by either an exception or an interrupt

So, not all exceptions and interrupts cause traps. Most notably, floating point exceptions and disabled interrupts do not trap.

If an exception or interrupt doesn't trap, the program counter does not change. So, there is no need to trace all exceptions/interrupts, just traps.

In this document, interrupts and exceptions are only traced when they cause traps to be taken.

Chapter 2

Branch Trace

Instruction delta tracing, also known as branch tracing, works by tracking execution from a known start address by sending information about the deltas taken by the program. Deltas are typically introduced by jump, call, return and branch type instructions, although interrupts and exceptions are also types of deltas.

Instruction delta tracing provides an efficient encoding of an instruction sequence by exploiting the deterministic way the processor behaves based on the program it is executing.

The approach relies on an offline copy of the program binary being available to the decoder, so it is generally unsuitable for either dynamic (self-modifying) programs or those where access to the program binary is prohibited.

While the program binary is sufficient, access to the assembly or higher-level source code will improve the ability of the decoder to present the decoded trace in the debugger by annotating the traced instructions with source code line numbers and labels, variable names etc.

This approach can be extended to cope with small sections of deterministically dynamic code by arranging for the decoder to request instruction memory from the target. Memory lookups generally lead to a prohibitive reduction in performance, although they are suitable for examining modest jump tables, such as the exception/interrupt vector pointers of an operating system which may be adjusted at boot up and when services are registered. Both static and dynamically linked programs can be traced using this approach. Statically linked programs are straightforward as they generally operate in a known address space, often mapping directly to physical memory. Dynamically linked programs require the debugger to keep track of memory allocation operations using either trace or stop-mode debugging.

2.1 Instruction delta trace concepts

2.1.1 Sequential instructions

For instruction set architectures such as RISC-V where all instructions are executed unconditionally or at least their execution can be determined based on the program binary, the instructions between the deltas are assumed to be executed sequentially. Consequently, there is no need to report them in the trace. The trace only needs to contain whether branches were taken or not, the addresses of taken indirect jumps, or other program counter discontinuities.

2.1.2 Uninferable PC discontinuities

An uninferable program counter discontinuity is a program counter change that can not be inferred from the program binary alone. For these cases, the instruction delta trace must include a destination address: the address of the next valid instruction.

Examples of this are indirect jumps, where the next instruction address is determined by the contents of a register rather than a constant embedded in the program binary.

2.1.3 Branches

A branch is an instruction where a jump is conditional on the value of a register or a flag. For a decoder to able to follow program flow, the trace must include whether a branch was taken or not.

For a direct branch, where the destination address is a constant that is encoded in the program binary, no further information is required. Direct branches are the only type of branch that is supported by the RISC-V ISA.

2.1.4 Interrupts and exceptions

Interrupts are a different type of delta that generally occur asynchronously to the program's execution rather than intentionally as a result of a specific instruction or event. Exceptions can be thought of in the same way, even though they can be typically linked back to a specific instruction address.

The decoder generally does not know where an interrupt occured in the instruction sequence, so the trace must report the address where normal program flow ceased, as well as give an indication of the asynchronous destination which may be as simple as reporting the exception type. When an interrupt or exception occurs, or the processor is halted, the final instruction executed beforehand must be traced. Following this, for an interrupt or exception, the next valid instruction address (the first of the interrupt or exception handler) must be traced in order to instruct the trace decoder to classify the instruction as an indirect jump even if it is not.

2.1.5 Synchronization

In order to make the trace robust there must be regular synchronization points within the trace. Synchronization is accomplished by sending a full valued instruction address (and potentially a context identifier). The decoder and debugger may also benefit from sending the reason for synchronizing. The frequency of synchronization is a trade-off between robustness and trace bandwidth.

The instruction trace encoder needs to synchronise fully:

- After a reset.

- When tracing starts.

- If the instruction is the first of an interrupt service routine or exception handler (hardware context change).

- After a prolonged period of time.

2.2 Optional and run-time configurable modes

An instruction trace encoder may support multiple tracing modes. To ensure that the decoder treats the incoming packets correctly, it needs to be informed of the current active configuration. The configuration is reported by a packet that is issued by the encoder whenever the encoder configuration is changed.

Here are common examples of such modes:

- delta address mode: program counter discontinuities are encoded as differences instead of absolute address values.

- full address mode: program counter discontinuities are encoded as absolute address values.

- implicit exception mode: the destination address of an exception (i.e. the address of the exception trap) is assumed to be known by the decoder, and thus not encoded in the trace.

- Sequentially inferable jump mode: The target of an indirect jump can be infered by considering the combined effect of two instructions.

- implicit return mode: the destination address of function call returns is derived from a call stack, and thus not encoded in the trace.

- branch prediction mode: branches that are predicted correctly by an encoder branch predictor (and an identical copy in the decoder) are not encoded as taken/non-taken, but as a more efficient branch count number.

- Jump target cache mode: Rather than reporting the address of an uninferable jump target, efficiency can be improved by caching recent jump targets, and reporting the cache entry index instead.

Modes may have associated parameters; see Table 7.1 for further details.

All modes are optional apart from delta address mode, which must be supported.

2.2.1   Delta address mode

Related parameters: None

In delta address mode, addresses are encoded as the difference between the actual address of the current instruction and the actual address of the instruction reported in the previous packet that contained an address. This differential encoding requires fewer bits than the full address, and thus results in more efficient trace compression.

2.2.2   Full address mode

Related parameters: None

In full address mode, all addresses in the trace are encoded as absolute addresses instead of in differential form. This kind of encoding is always less efficient, but it can be a useful debugging aid for software decoder developers.

2.2.3   Implicit exception mode

Related parameters: None

The RISC-V Privileged ISA specification stores exception handler base addresses in the **utvec/stvec/mvec** CSR registers. In some RISC-V implementations, the lower address bits are stored in the **ucause/scause/mcause** CSR registers.

By default, both the ***vec** and ***cause** values are reported when an exception or interrupt occurs.

The implicit exception mode omits ***vec** (the trap handler address), from the trace and thus improves efficiency.

This mode can only be used if the decoder can infer the address of the trap handler from just the exception cause.

2.2.4   Sequentially inferable jump mode

Related parameters: *sijump_p*.

By default, the target of an indirect jump is always considered an uninferable PC discontinuity. However, if the register that specifies the jump target was loaded with a constant then it can be considered inferable under some circumstances. The hart must identify jumps with sequentially inferable targets and provide this information separately to the encoder. The final decision as to

whether to treat the jump as inferable or not must be made by the encoder. Both the constant load and the jump must be traced in order for the decoder to be able to infer the jump target.

2.2.5 Implicit return mode

Related parameters: *call_counter_size_p*, *return_stack_size_p*.

Although a function return is usually an indirect jump, well behaved programs return to the point in the program from which the function was called using a standard calling convention. For those programs, it is possible to determine the execution path without being explicitly notified of the destination address of the return. The implicit return mode can result in very significant improvements in trace encoder efficiency.

Returns can only be treated as inferable if the associated call has already been reported in an earlier packet. The encoder must ensure that this is the case. This can be accomplished by utilizing a counter to keep track of the number of nested calls being traced. The counter increments on calls (but not tail calls), and decrements on returns (see Section 3.1.1 for definitions). The counter will not over or underflow, and is reset to 0 whenever a synchronization packet is sent. Returns will be treated as inferable and will not generate a trace packet if the count is non-zero (i.e. the associated call was already reported in an earlier packet).

Such a scheme is low cost, and will work as long as programs are "well behaved". The encoder does not check that the return address is actually that of the instruction following the associated call. As such, any program that modifies return addresses cannot be traced using this mode with this minimal implementation.

Alternatively, the encoder can maintain a stack of expected return addresses, and only treat a return as inferable if the actual return address matches the prediction. This is fully robust for all programs, but is more expensive to implement. In this case, if a return address does not match the prediction, it must be reported explicitly via a packet, along with the number of return addresses currently on the stack. This ensures that the decoder can determine which return is being reported.

2.2.6 Branch prediction mode

Related parameters: *bpred_size_p*.

Without branch prediction, the outcome of each executed branch is stored in a branch map: a bit vector in which the taken/non-taken status of each branch is stored in chronological order.

While this encoding is efficient, at 1 bit per branch, there are some cases where this can still result in a relatively large volume of trace packets. For example:

- Executing tight loops of code containing no uninferable jumps. Each iteration of the loop will add a bit to the branch map;

- Sitting in an idle loop waiting for an interrupt. This produces large amounts of trace when nothing of any interest is actually happening!

- Breakpoints, which in some implementations also spin in an idle loop.

A significant coding efficiency can be obtained by the addition of a branch predictor in the encoder. To keep the encoder and decoder synchronized, a predictor with identical behavior will need to be implemented in the decoder software.

The predictor shall comprise a lookup table of $2^{bpred\_size\_p}$ entries. Each entry is indexed by bits N:1 of the instruction address (or N+1:2 if compressed instructions aren't supported), and each contains a 2-bit prediction state:

- 00: predict not taken, transition to 01 if prediction fails;

- 01: predict not taken, transition to 00 if prediction succeeds, else 11;

- 11: predict taken, transition to 10 if prediction fails;

- 10: predict taken, transition to 11 if prediction succeeds, else 00.

The MSB represents the predicted outcome, the LSB the most recent actual outcome. The prediction must fail twice for the predicted value to change.

The lookup table entries are initialized to 01 when a synchronization packet is sent.

Other predictors, such as the gShare predictor (see Hennessy & Patterson), should be considered. Some further experimentation is needed to determine the benefits of different lookup table sizes and predictor algorithms.

2.2.7   Jump target cache mode

Related parameters: *cache_size_p*.

By default, the target address of an uninferable jump is output in the trace, usually in differential form. If the same function is called repeatedly, (for example, in a loop), the same address will be output repeatedly.

An efficiency gain can be obtained by the addition of a jump target cache to the encoder. To keep the encoder and decoder synchronized, a cache with identical behavior will need to be implemented in the decoder software. Even a small cache can provide significant improvement.

The cache shall comprise $2^{cache\_size\_p}$ entries. It will be direct mapped, with each entry indexed by bits N:1 of the instruction address (or N+1:2 if compressed instructions aren't supported).

Each uninferable jump target is first compared with the entry at its index in the cache. If it is found in the cache, the index number is traced rather than the target address. If it is not found in the cache, the entry at that index is replaced.

Chapter 3

Ingress Port

3.1 Interface requirements

This section describes in general terms the information which must be passed from the RISC-V hart to the trace encoder, and distinguishes between what is mandatory, and what is optional.

The following information is mandatory:

- The number of instructions that are being retired;

- Whether there has been an exception or interrupt, and if so the cause (from the *ucause/scause/mcause* CSR) and trap value (from the *utval/stval/mtval* CSR);

- The current privilege level of the RISC-V hart;

- The *instruction_type* of retired instructions for:

    - Jumps with a target that cannot be inferred from the source code;
    - Taken branches;
    - Return from exception or interrupt (**ret* instructions).

- The *instruction_address* for:

    - Jumps with a target that *cannot* be inferred from the source code;
    - Taken branches;
    - The instruction executed immediately after a jump or taken branch (also referred to as the target or destination of the jump or taken branch);
    - The last instruction executed before an exception or interrupt;
    - Exceptions;
    - The first instruction executed following an exception or interrupt;
    - The last instruction executed before a privilege change;
    - The first instruction executed following a privilege change;

  – The first and last instruction being retired.

- The number of nontaken branches being retired.

The following information is optional:

- Context information:

  – The context and/or hart ID;
  – The type of action to take when context changes.

- The *instruction_type* of instructions for:

  – Calls with a target that *cannot* be inferred from the source code;
  – Calls with a target that *can* be inferred from the source code;
  – Tail-calls with a target that *cannot* be inferred from the source code;
  – Tail-calls with a target that *can* be inferred from the source code;
  – Returns with a target that *cannot* be inferred from the source code;
  – Returns with a target that *can* be inferred from the source code;
  – Co-routine swap;
  – Jumps which don't fit any of the above classifications with a target that *cannot* be inferred from the source code;
  – Jumps which don't fit any of the above classifications with a target that *can* be inferred from the source code;
  – Nontaken branches.

- If context is supported then the *instruction_address* for:

  – The last instruction executed before a context change;
  – The first instruction executed following a context change.

- Whether jump targets are sequentially inferable or not.

The mandatory information is the bare-minimum required to implement the branch trace algorithm outlined in Chapter 5. The optional information facilitates alternative or improved trace algorithms:

- Implicit return mode (see Section 2.2.5) requires the encoder to keep track of the number of nested function calls, and to do this it must be aware of all calls and returns regardless of whether the target can be inferred or not;

- A simpler algorithm useful for basic code profiling would only report function calls and returns, again regardless of whether the target can be inferred or not;

- Branch prediction techniques can be used to further improve the encoder efficiency, particularly for loops (see Section 2.2.6). This requires the encoder to be aware of the address of all branches, whether they are taken or not.

- Uninferable jumps can be treated as inferable (which don't need to be reported in the trace output) if both the jump and the preceding instruction which loads the target into a register have been traced.

3.1.1   Jump classification and target inference

Jumps are classified as *inferable*, or *uninferable*. An *inferable* jump has a target which can be deduced from the binary executable or representation thereof (e.g. ELF). For the purposes of this specification, the following strict definition applies:

If the target of a jump is supplied via a constant embedded within the jump opcode, it is classified as *inferable*. Jumps which are not *inferable* are by definition *uninferable*.

However, there are some jump targets which can still be deduced from the binary executable by considering pairs of instructions even though by the above definition they are classified as uninferable. Specifically, jump targets that are supplied via

- an **lui** or **c.lui** (a register which contains a constant), or

- an **auipc** (a register which contains a constant offset from the PC).

Such jump targets are classified as *sequentially inferable* if the pair of instructions are executed consecutively (i.e. the **auipc**, **lui** or **c.lui** immediately precedes the jump). Note: the restriction that the instructions are executed consecutively is necessary in order to minimize the additional signalling needed between the hart and the encoder, and should have a minimal impact on trace efficiency as it is anticipated that consecutive execution will be the norm. Support for sequentially inferable jumps is optional.

Jumps may optionally be further classified according to the recommended calling convention:

- *Calls*:

    - **jal** x1;
    - **jal** x5;
    - **jalr** x1, rs where rs != x1;
    - **jalr** x5, rs where rs != x5;
    - **c.jalr** rs1.

- *Tail-calls*:

    - **jalr** x0, rs where rs != x1 and rs != x5;
    - **c.jr** rs1 where rs1 != x1 and rs1 != x5.

- *Returns*:

    - **jalr** x0, rs where rs == x1 or rs == x5;
    - **c.jr** rs1 where rs1 == x1 or rs1 == x5.

- *Co-routine swap*:

    - **jalr** x1, x1;

    – *jalr* x5, x5.

- *Other*:

    – *jal* rd where rd != x1 and rd != x5;
    – *jalr* rd, rs where rd != x0 and rd != x1 and rd != x5.

3.2   Instruction interface

This section describes the interface between a RISC-V hart and the trace encoder that conveys the information described in the previous section.

Tables 3.1, 3.2 and 3.3 list the signals in the interface designed to efficiently support retirement of multiple instructions per cycle. The following discussion describes the multiple-retirement behavior. However, for harts that can only retire one instruction at a time, the signalling can be simplified, and this is discussed subsequently in Section 3.2.1.

The information presented on the ingress port represents a contiguous block of instructions starting at **iaddr**, all of which retired in the same cycle. Note if **itype** is 1 or 2 (indicating an exception or an interrupt), the number of instructions retired may be zero. **cause** and **tval** are only defined if **itype** is 1 or 2. If **iretire**=0 and **itype**=0, the values of all other signals are undefined.

iretire contains the number of half-words represented by instructions retired in this block, and **ilastsize** the size of the last instruction. Half-words rather than instruction count enables the encoder to easily compute the address of the last instruction in the block without having access to the size of every instruction in the block.

If address translation is enabled, **iaddr** is a virtual address, else it is a physical address. Virtual addresses narrower than *iaddress_width_p* bits must be sign-extended to make computation of differential addresses easier, and physical addresses narrower than *iaddress_width_p* bits must be zero-extended.

Cores can retire multiple non-taken branches per clock cycle, indicated via **ntkn**. However, a consequence of this is that the encoder will be unaware of the addresses of some non-taken branches, which will prevent the use of a branch predictor to improve compression (see Section 2.2.6). For harts that can only retire a maximum of one non-taken branch per clock cycle, **ntkn** can be omitted, provided all non-taken branches are indicated via **itype**. The number of non-taken branches is **ntkn** if **ntkn** is non-zero, or 1 if **itype** = 4 and **ntkn** is zero. In other words, if for example **ntkn** is 2 and **itype** = 4, the encoder will interpret this as 2 non-taken branches, not 3.

For harts that can retire a maximum of N taken branches per clock cycle, the signal group (**iretire**, **itype**, **ntkn** (if present), **ilastsize**, **iaddr**) must be replicated N times. Signal group 0 represents information about the oldest instruction block, and group N-1 represents the newest instruction block. The interface supports no more than one privilege, context, exception or interrupt per cycle and so **priv**, **context**, **context_type**, **cause** and **tval** are not replicated. Furthermore, **itype** can only take the value 1 or 2 in one of the signal groups, and this must be the newest valid group (i.e. **iretire** and **itype** must be zero for higher numbered groups). If fewer than N taken branches are retired in a cycle, then lower numbered groups must be used first. For example, if there is one

Table 3.1: Instruction interface signals - common

Signal	Function
itype[*itype_width_p*-1:0]	Termination type of the instruction block (see Section 3.1.1 for definitions of codes 6 - 15): 0: Final instruction in the block is none of the other named **itype** codes; 1: Exception. An exception occurred following the final retired instruction in the block; 2: Interrupt. An interrupt occurred following the final retired instruction in the block; 3: Exception return; 4: Nontaken branch; 5: Taken branch; 6: reserved; 7: Co-routine swap; 8: Uninferable call; 9: Inferrable call; 10: Uninferable tail-call; 11: Inferrable tail-call; 12: Uninferable return; 13: Inferrable return; 14: Other uninferable jump; 15: Other inferable jump.
cause[*ecause_width_p*-1:0]	Exception or interrupt cause (**ucause/scause/mcause**), Ignored unless **itype**=1 or 2.
tval[*iaddress_width_p*-1:0]	The associated trap value, e.g. the faulting virtual address for address exceptions, as would be written to the **utval/stval/mtval** CSR. Future optional extensions may define **tval** to provide ancillary information in cases where it currently supplies zero Ignored unless **itype**=1 or 2.
priv[*privilege_width_p*-1:0]	Privilege level for all instructions in this block.
context[*context_width_p*-1:0]	Context and/or hart ID for all instructions in this block.
iaddr[*iaddress_width_p*-1:0]	The address of the 1st instruction retired in this block. Invalid if **iretire**=0
context_type[*context_type_width_p*-1:0]	Behavior type of **context** 0: Context change with discontinuity; 1: Precise context change; 2: Imprecise context change; 3: Notification.

Table 3.2: Instruction interface signals - multiple retirement

Signal	Function
iretire[*iretire_width_p*-1:0]	Number of halfwords represented by instructions retired in this block.
ilastsize[*ilastsize_width_p*-1:0]	The size of the last retired instruction is $2^{\textbf{ilastsize}}$ half-words.

Table 3.3: Instruction interface signals - multiple non-taken branches

Signal	Function
ntkn[*ntkn_width_p*-1:0]	Number of nontaken branches in this block.

taken branch, use only group 0, if there are two taken branches, instructions up to the 1st taken branch must be reported in group 0 and instructions up to the 2nd taken branch must be reported in group 1 and so on.

sijump is optional and may be omitted if the hart does not implement the logic to detect sequentially inferable jumps. If the encoder offers an **sijump** input it must also provide a parameter to indicate whether the input is connected to a hart that implements this capability, or tied off. This is to ensure the decoder can be made aware of the hart's capability. Enabling sequentially inferable jump mode in the encoder and decoder when the hart does not support it will prevent correct reconstruction by the decoder.

The **context** field can be used to convey any additional information to the decoder. For example:

- The hart ID;

- The software thread ID;

- It could be used to convey the values of CSRs to the decoder by setting **context** to the CSR number and value when a CSR is written.

Table 3.6 specifies the actions for the various **context_type** values.

3.2.1   Simplifications for single-retirement

For harts that can only retire one instruction at a time, the interface can be simplified to the signals listed in tables 3.1 and 3.5. The simplifications can be summarized as follows:

- As the number of instructions that are retired in a block is only 0 or 1, the encoder does not need information to enable it to deduce the address of the last instruction retired (it is the same as the 1st and only instruction retired). So **ilastsize** is not necessary, and **iretire** simply indicates whether an instruction retired or not;

- As the number of non-taken branches retired is never more than 1, and can always be indicated via **itype**, **ntkn** is not necessary.

Table 3.4: Instruction interface signals - sequentially inferable jumps

Signal	Function
sijump	If **itype** indicates that this block ends with an uninferable discontinuity, setting this signal to 1 indicates that it is sequentially inferable and may be treated as inferable by the encoder if the preceding *auipc*, *lui* or *c.lui* has been traced. Must be zero for **itype** codes other than 8, 10, 12 or 14.

Table 3.5: Instruction interface signals - single retirement

Signal	Function
iretire[0:0]	Number of instructions retired in this block (0 or 1).

Table 3.6: Call/return **context_type** values and corresponding actions

Type	Value	Actions
Context change with discontinuity	0	An example would be a change of hart. Need to report the last instruction executed on the previous context, as well as the 1st on the new context. Treated the same as an exception.
Precise context change	1	Need to output the address of the 1st instruction, and the new context. If there were unreported branches beforehand, these need to be output first. Treated the same as a privilege change.
Imprecise context change	2	An example would be a SW thread change. Report the new context value at the earliest convenient opportunity. It is reported without any address information, and the assumption is that the precise point of context change can be deduced from the source code (e.g. a CSR write).
Notification	3	An example would be a *Trace-Notify* from the hart's trigger unit (see section 3.2.4). Need to output the address of the last instruction in the block. The context itself is not output.

The parameter *retires_p* which indicates to the encoder the maximum number of instructions that can be retired per cycle can be used by an encoder capable of supporting single or multiple retirement to select the appropriate interpretation of **iretire**. **ilastsize** and **ntkn** encoder inputs must be tied low when attached to a single-retirement hart that does not provide these outputs.

3.2.2   Alternative multiple-retirement interface configurations

For a hart that can retire multiple instructions per cycle, but no more than one taken branch, the preferred solution is to use one of each of the signals from tables 3.1, 3.2 and optionally 3.3. However, an alternative approach would be to provide explicit details of every instruction retired by using N sets of the signal group (**iretire**, **itype**) from tables 3.1 and 3.5 with the groups detailing one instruction each (replicating the single retirement example N times).

3.2.3   Optional sideband signals

Optional sideband signals may be included to provide additional functionality, as described in tables 3.7 and 3.8.

Note, any user defined information that needs to be output by the encoder will need to be applied via the **context** input.

Table 3.7: Optional sideband encoder ingress signals

Signal	Function
impdef[*impdef_width_p*-1:0]	Implementation defined sideband signals. A typical use for these would be for filtering (see Chapter 4.
trigger[1:0]	A pulse on bit 0 will cause the encoder to start tracing, and continue until further notice, subject to other filtering criteria also being met. A pulse on bit 1 will cause the encoder to stop tracing until further notice. See section 3.2.4).
halted	Core is halted. Upon assertion, the encoder will output a packet to report the address of the last instruction retired before halting, followed by a support packet to indicate that tracing has stopped. Upon deassertion, the encoder will start tracing again, commencing with a synchronization packet.
reset	Core is in reset. Provided the encoder is in a different reset domain to the hart, this allows the encoder to indicate that tracing has ended on entry to reset, and restarted on exit. Behavior is as described above for halt.

Table 3.8: Optional sideband encoder egress signals

Signal	Function
stall	Stall request to hart. Some applications may require lossless trace, which can be achieved by using this signal to stall the hart if the trace encoder is unable to output a trace packet (for example due to back-pressure from the packet transport infrastructure).

3.2.4 Using trigger outputs from the Debug Module

The debug module of the RISC-V hart may have a trigger unit. This defines a match control register (*mcontrol*) containing a 4-bit **action** field, and reserves codes 2 - 5 of this field for trace use. These action codes are hereby defined as shown in table 3.9. The table also includes details of how these actions should be comunicated to the trace encoder.

Table 3.9: Debug Module trigger support (*mcontrol* **action**)

Value	Description
2	*Trace-on.* Generate a pulse on an output from the hart. This should be connected to **trigger[0]** if the encoder provides it.
3	*Trace-off.* Generate a pulse on an output from the hart. This should be connected to **trigger[1]** if the encoder provides it.
4	*Trace-notify.* Set the **context_type** encoder input to 'Notification' for the block of instructions that ends with the instruction that caused the triggered the notification.

Trace-on and Trace-off actions provide a means for the hart to control when tracing occurs. Trace-notify provides means to ensure that a specified instruction is explicitly reported. This capability is sometimes known as a watchpoint.

3.2.5 Example retirement sequences

Table 3.10: Example 1 : 9 Instructions retired over three cycles, 2 branches

Retired	Instruction Trace Block
1000: *divuw* 1004: *add* 1008: *or* 100C: *c.jalr*	**iretire=7, iaddr=0x1000, ntkn=0, itype=8**
0940: *addi* 0944: *c.beq* 0946: *c.bnez*	**iretire=4, iaddr=0x0940, ntkn=1, itype=5**
0988: *lbu* 098C: *csrrw*	**iretire=4, iaddr=0x0988, ntkn=0, itype=0**

Chapter 4

Filtering

Filtering provides a mechanism to control whether the encoder should produce trace. For example, it may be desirable to trace:

- When the instruction address is within a particular range;

- Starting from one instruction address and continuing until a second instruction address;

- For one or more specified privilege levels;

- For a particular context or range of contexts;

- Exception and/or interrupt handlers for specified exception causes or with particular **tval** values;

- Based on values applied to the **impdef** or **trigger** signals;

- For a fixed period of time

- etc.

How this is accomplished is implementation specific.

One suggested implementation provides:

- Comparators offering a range of arithmetic options ($<$, $>$, $=$, $!=$, etc) for **iaddress**, **context** and **tval** inputs;

- Multiple choice selection for **priv** and **cause** inputs;

- Masked matching for **interrupt** and **impdef** inputs;

- The ability to enable tracing when **trigger[0]** is asserted, and continue tracing until **trigger[0]** is asserted (see Section 3.2.4).

Chapter 5

Example Algorithm

An example algorithm for compressed branch trace is given in figure 5.1. In the diagram, the following terms are used:

- *te_inst.* The name of the packet type emitted by the encoder (see Chapter 6);

- *inst.* Abbreviation for 'instruction';

- *updiscon.* Uninferable PC discontinuity. This identifies an instruction that causes the program counter to be changed by an amount that cannot be predicted from the source code alone (**itype** values 8, 10, 12 or 14);

- *Qualified?* An instruction that meets the filtering criteria is qualified, and will be traced;

- *Branch?* Is the instruction a branch or not (**itype** values 4 or 5, or a non-zero **ntkn**);

- *branch map.* A vector where each bit represents the outcome of a branch. A 0 indicates the branch was taken, a 1 indicates that it was not;

- *e_ccd.* An exception has been signalled, or context has changed and should be treated as an uninferable PC discontinuity (see Table 3.6);

- *ppch.* Privilege has changed, or context has changed and needs to be reported precisely (see Table 3.6);

- *ppch_br.* As above, but branch map not empty;

- *er_ccdn.* Instruction retirement and exception signalled on the same cycle, or context has changed and should be treated as an uninferable PC discontinuity, or context notification (see Table 3.6);

- *exc_only.* Exception signalled without simultaneous retirement;

- *cci.* context change that can be reported imprecisely (see Table 3.6);

- *resync count.* A counter used to keep track of when it is necessary to send a synchronization packet (see Section 5.2);

- *max_resync.* The resync counter value that schedules a synchronization packet (see Section 5.2);

- *resync_br.* The resync counter has reached the maximum value and there are entries in the branch map that have not yet been output (see Section 5.2).

Figure 5.1 shows instruction by instruction behavior, as would be seen in a single-retirement system only. Whilst the ingress port allows the RISC-V hart to provide information on multiple retiring instructions simultaneously, the resultant packet sequence generated by the encoder must be the same as if retiring one instruction at a time.

A 3-stage pipeline within the encoder is assumed, such that the encoder has visibility of the current, previous and next instructions. All packets are generated using information relating to the current instruction. The orange diamonds indicate decisions based on the previous (or last) instruction, the green diamond indicates a decision based on the next instruction, and all other diamonds are based on the current instruction.

Additionally, the encoder can generate one further packet type, not shown on the diagram for clarity. The *support* packet (format 3, subformat 3 - see Chapter 6) is sent when:

- The encoder is enabled or disabled, or its configuration is changed, to inform the decoder of the operating mode of the encoder

- After the last qualified instruction has been traced, to inform the decoder that tracing has stopped;

- If trace packets are lost (for example if the buffer into which packets are being written fills up. In this situation, the 1st packet loaded into the buffer when space next becomes available should be a *support* packet. Following this, tracing will resume with a sync packet.

Note: if the **halted** or **reset** sideband signals are asserted (see Table 3.7) the encoder will behave as if it has received an unqualified instruction (output *te_inst* reporting the address of the last instruction, followed by *te_support*);

5.1   Format selection

In all cases but one, the packet format is determined only by a 'yes' outcome from the associated decision. The choice between formats 1 or 2 for the case in the middle of the diagram needs further explanation.

If there are no branches that need to be reported, packet format 2 is used.

If there are branches to report, format 1 is used. If branch prediction is supported and is enabled, then the encoder may choose whether to output a full branch map, or a count of correctly predicted branches.

Figure 5.1: Instruction delta trace algorithm

Packet formats 1 and 2 are organized so that the address is usually the final field. Minimizing the number of bits required to represent the address reduces the total packet size and significantly improves efficiency. See Chapter 6.

5.2   Resynchronisation

Per Section 2.1.5, a format 3 synchronisation packet must be output after "a prolonged period of time". The exact mechanism for determining this is not specified, but options might be to count the number of *te_inst* packets emitted, or the number of clock cycles elapsed, since the last synchronization message was sent.

When the resync is required, the primary objective is to output a format 3 packet, so that the decoder can start tracing from that point without needing any of the history. However, if the decoder is already synced, then it is also required that it can continue to follow the execution path up to and through the format 3 packet seamlessly. As such, before outputting a format 3 packet, it is necessary to output a format 1 packet for the preceding instruction if there are any unreported branches (because format 3 does not contain a branch map). The format 3 will be sent if the resync timer has been exceeded. On the cycle before this (when the resync timer value has been exactly reached), a format 1 will be generated if the branch map is not empty.

Chapter 6

Trace Encoder Output Packets

The bulk of this section describes the payload of packets output from the Trace Encoder. The infrastructure used to transport these packets is outside the scope of this document, and as such the manner in which packets are encapsulated for transport is not specified. However, the following information must be provided to the encapsulator:

- The packet type;

- The packet length, in bytes;

- The packet payload.

Two example transport schemes are the UltraSoC Messaging Infrastructure, and the Arm Trace Bus. Figure 6.1 shows the encapsulation used for the UltraSoC infrastructure:

- The header byte contains a 5-bit field specifying the payload length in bytes, a 2-bit field indicating the "flow" (destination routing indicator), and a bit to indicate whether an optional 16-bit timestamp is present;

- The index field indicates the source of the packet. The number of bits is system dependent, And the initial value emitted by the trace encoder is zero (it gets adjusted as it propagates through the infrastructure);

- An optional 2-byte timestamp;

- The packet payload.

Figure 6.1: Example encapsulated packet format

Alternatively, for ATB, the source of the packet is indicated by the **ATID** bus field, and there is no equivalent of "flow", so an example encapsulation might be:

27

- A 5-bit field specifying the payload length in bytes

- A bit to indicate whether an optional 16-bit timestamp is present;

- An optional 2-byte timestamp;

- The packet payload.

It may be desirable for packets to start aligned to an ATB word, in which the **ATBYTES** bus field in the last beat of a packet can be used to indicate the number of valid bytes.

The remainder of this section describes the contents of the payload portion which should be independent of the infrastructure. In each table, the fields are listed in transmission order: first field in the table is transmitted first, and multi-bit fields are transmitted LSB first.

This packet payload format is used to output encoded instruction trace. Three different formats are used according to the needs of the encoding algorithm. The following tables show the format of the payload - i.e. excluding any encapsulation.

In order to achieve best performance, actual packet lengths may be adjusted using 'sign based compression'. At the very minimum this should be applied to the address field of format 1 and 2 packets, but ideally will be applied to the whole packet, regardless of format. This technique eliminates identical bits from the most significant end of the packet, and adjusts the length of the packet accordingly. A decoder receiving this shortened packet can reconstruct the original full-length packet by sign-extending from the most significant received bit.

Where the payload length given in the following tables, or after applying sign-based compression, is not a multiple of whole bytes in length, the payload must be sign-extended to the nearest byte boundary.

Whilst offering maximum encoding efficiency, variable length packets can present some challenges, specifically in terms of identifying where the boundaries between packets occur either when packed packets are written to memory, or when packets are streamed offchip via a communications channel. Two potential solutions to this are as follows:

- If the maximum packet payload length is 2^N-1 (for example, if N is 5, then the maximum length is 31 bytes), and the minimum packet payload length is 1, then a sequence of at least 2^N zero bytes cannot occur within a packet payload, and therefore the first non-zero byte seen after a sequence of at least 2^N zero bytes must be the first byte of a packet. This approach can be used for alignment in either memory or a data stream;

- An alternative approach suitable for packets written to memory is to divide memory into blocks of M bytes (e.g. 1kbyte blocks), and write packets to memory such that the first byte in every block is always the first byte of a packet. This means packets cannot span block boundaries, and so zero bytes must be used to pad between the end of the last message in a block and the block boundary.

6.1   Format 3 packets

Format 3 packets are used for synchronization, reporting context and supporting information. There are 4 sub-formats.

Throughout this document, the term "synchronization packet" is used. This refers specifically to format 3, subformat 0 and subformat 1 packets.

6.2   Format 3 subformat 0 - Synchronisation

This packet contains all the information the decoder needs to fully identify an instruction. It is sent for the first traced instruction (unless that instruction also happens to be a the first in an exception handler), and when resynchronization has been scheduled by expiry of the resynchronisation timer.

Table 6.1: Packet format 3, subformat 0

Field name	Bits	Description
format	2	11 (sync): synchronisation
subformat	2	00 (start): Start of tracing, or resync
context	*context_width_p*, or 0 if *nocontext_p* is 1	The instruction context
privilege	*privilege_width_p*	The privilege level of the reported instruction
address	*iaddress_width_p* - *iaddress_lsb_p*	Full instruction address. Address alignment is determined by *iaddress_lsb_p* Address must be left shifted in order to recreate original byte address.
branch	1	If the address points to a branch instruction, the branch is not taken if the value of this bit is different from the MSB of **address**. Set to the same value as the MSB of **address** if the branch is taken or the instruction is not a branch.

6.2.1   Format 3 branch field

This bit indicates the taken/not taken status in the case where the reported address points to a branch instruction. Overall efficiency would be slightly improved if this bit was removed, and the branch status was instead "carried over" and reported in the next *te_inst* packet. This was considered, but there are several pathological cases where this approach fails. Consider for example the situation where the first instruction that matches the filtering criteria is a branch, and this is then followed immediately by an exception. This results in format 3 packets being generated on two consecutive cycles. The second packet does not contain a branch map, so there is no way to report the branch status of the 1st branch, apart from by inserting a format 1 packet in between. There are two issues with this:

- It would require the generation of 2 packets on the same cycle, which adds significant additional complexity to the encoder;

- It would complicate the algorithm shown in 5.1.

This bit is encoded so that most of the time it will take the same value as the MSB of the preceding field, and will therefore compress away, in order to minimize the efficiency impact. Branches are unlikely to be reported using a format 3 packet apart from if the 1st traced instruction is a branch, or if the instruction reported when the resync timer expires is a branch.

6.3   Format 3 subformat 1 - Exception

This packet also contains all the information the decoder needs to fully identify an instruction. It is sent following an exception, and as well as reporting the address of the exception handler, it also includes the exception cause and the address of the faulted instruction.

If the implicit exception mode is enabled (see section 2.2.3), the address is omited.

Table 6.2: Packet format 3, subformat 1

Field name	Bits	Description
format	2	11 (sync): synchronisation
subformat	2	01 (exception): Exception cause and trap handler address.
context	*context_width_p*, or 0 if *nocontext_p* is 1	The instruction context.
privilege	*privilege_width_p*	The privilege level of the reported instruction.
address	*iaddress_width_p* - *iaddress_lsb_p*	Full instruction address. Address alignment is determined by *iaddress_lsb_p* Address must be left shifted in order to recreate original byte address.
ecause	*ecause_width_p*	Exception cause.
interrupt	1	Interrupt.
tvalepc	*iaddress_width_p*	Exception address if **ecause** is 2 and **interrupt** is 0 (illegal instruction exception), or trap value otherwise.
branch	1	If the address points to a branch instruction, the branch is not taken if the value of this bit is different from the MSB of **tvalepc**. Set to the same value as the MSB of **tvalepc** if the branch is taken or the instruction is not a branch.

6.3.1   Format 3 tvalepc field

This field reports the address of illegal instructions, or the trap value otherwise. This ensures that the address of the faulting instruction is reported for all required cases. The trap value is set to the address of the faulting instruction for hardware breakpoints, access or page faults and instructions, loads or stores that are mis-aligned, but not for illegal instructions (for which it is set to the opcode).

6.4 Format 3 subformat 2 - Context

This packet contains only the context, and is output when the context changes and can be reported imprecisely (see Table 3.6).

Table 6.3: Packet format 3, subformat 2

Field name	Bits	Description
format	2	11 (sync): synchronisation
subformat	2	10 (context): Context change
context	*context_width_p*	The instruction context.
privilege	*privilege_width_p*	The privilege level of the new context.

6.5 Format 3 subformat 3 - Support

This packet provides supporting information to aid the decoder. It is issued when

- Trace is enabled or disabled;

- The operating mode changes;

- One or more trace packets cannot be sent (for example, due back-pressure from the packet transport infrastructure).

The **options** field is a placeholder that must be replaced by an implementation specific set of individual bits - one for each of the optional modes supported by the encoder.

6.5.1 Format 3 subformat 3 qual_status field

When tracing ends, the encoder reports the address of the last traced instruction, and follows this with a format 3, subformat 3 (supporting information) packet. Two codes are provided for indicating that tracing has ended: **ended_rep** and **ended_upd**. This relates to exactly the same ambiguous case described in detail in section 6.6.2, and in principle, the mechanism described in that section can be used to disambiguate when the last traced instruction is at looplabel + 4. However, that mechanism relies on knowing when creating the format 1/2 packet, that a format 3 packet will be generated from the next instruction. This is possible because the encoding algorithm uses a 3-stage pipe with access to the previous, current and next instructions. However, decoding that the next instruction is a privilege change or exception is straightforward, but determining whether the next instruction meets the filtering criteria is much more involved, and this information won't typically be available, at least not without adding an additional pipeline stage, which is expensive. This means a different mechanism is required, and that is provided by having two codes to indicate that tracing has ended:

- **ended_rep** indicates that the preceding packet would not have been issued if tracing hadn't ended, which means that tracing stopped after executing looplabel in the 1st loop iteration;

Table 6.4: Packet format 3, subformat 3

Field name	Bits	Description
format	2	11 (sync): synchronisation
subformat	2	11 (support): Supporting information for the decoder
enable	1	Indicates if the encoder is enabled
encoder_mode	N	Identifies trace algorithm Details and number of bits implementation dependent. Currently Branch trace is the only mode defined, indicated by the value 0.
qual_status	2	Indicates qualification status 00 (no_change): No change to filter qualification 01 (ended_rep): Qualification ended, preceding **te_inst** sent explicitly to indicate last qualification instruction 10: (trace_lost): One or more packets lost. 11 : (ended_upd): Qualification ended, preceding **te_inst** would have been sent anyway due to an up-discon, even if it wasn't the last qualified instruction)
options	N	Values of all run-time configuration bits Number of bits and definitions implementation dependent. Examples might be - 'sequentially infered jumps' Don't report the targets of sequentially inferable jumps - 'implicit return' Don't report function return addresses - 'implicit exception' Exclude address from format 3, sub-format 1 *te_inst* packets if trap vector can be determined from *ecause field* - 'branch prediction' Branch predictor enabled - 'jump target cache' Jump target cache enabled - 'full address' Always output full addresses (SW debug option)

- **ended_upd** indicates that the preceding packet would have been issued anyway because of an uninferable PC discontinuity, which means that tracing stopped after executing looplabel in the 2nd loop iteration;

If the encoder implementation does have early access to the filtering results, and the designer chooses to use the **updiscon** bit when the last qualified instruction is also the instruction following an uninferable PC discontinuity, loss of qualification should always be indicated using **ended_rep**.

6.6   Format 2 packets

This packet contains only an instruction address, and is used when the address of an instruction must be reported, and there is no unreported branch information. The address is in differential format unless full address mode is enabled (see section 2.2.2).

Table 6.5: Packet format 2

Field name	Bits	Description
format	2	10 (addr-only): differential address and no branch information
address	*iaddress_width_p - iaddress_lsb_p*	Differential instruction address.
notify	1	If the value of this bit is different from the MSB of **address**, it indicates that this packet is reporting an instruction that is not the target of an uninferable discontinuity because a notification was requested via **context_type** (see table 3.1).
updiscon	1	If the value of this bit is different from **notify**, it indicates that this packet is reporting the instruction following an uninferable discontinuity and is also the instruction before an exception, privilege change or resync (i.e. it will be followed immediately by a format 3 *te_inst*).
irfail	1	If the value of this bit is different from **updiscon**, it indicates that this packet is reporting the instruction following a return because its address differs from the predicted return address at the top of the implicit_return return address stack.
irdepth	*return_stack_size_p*	If the value of **irfail** is different from **updiscon**, this field indicates the number of entries on the return address stack (i.e. the entry number of the return that failed). If **irfail** is the same value as **updiscon**, all bits in this field will also be the same value as **updiscon**.

6.6.1   Format 2 notify field

This bit is encoded so that most of the time it will take the same value as the MSB of the **address** field, and will therefore compress away, having no impact on the encoding efficiency. It is required in order to cover the case where an address is reported as a result of a notification request, signalled by setting the **context_type** input to 11.

6.6.2   Format 2 notify and updiscon fields

These bits are encoded so that most of the time they will compress away, having no impact on efficiency, by taking on the same value as the preceding bit in the packet (**notify** is normally the same value as the MSB of the **address** field, and **updiscon** is normally the same value as **notify**). They are required in order to cover a pathological case where otherwise the decoding software would not be able to reconstruct the program execution unambiguously. Consider the following code fragment:

looplabel - 4: ***opcode A***
looplabel   : ***opcode B***
looplabel + 4: ***opcode C***
:
looplabel + N: ***JALR*** # Jump to looplabel

This is a loop with an indirect jump back to the next iteration. This is an uninferable discontinuity, and will be reported via a format 1 or 2 packet. Note however that the initial entry into the loop is fall-through from the instruction at looplabel - 4, and will not be reported explicitly. This means that when reconstructing the execution path of the program, the looplabel address is encountered twice. On first glance, it appears that the decoder can determine when it reaches the loop label for the 1st time that this is not the end of execution, because the preceding instruction was not one that can cause an uninferable discontinuity. It can therefore continue reconstructing the execution path until it reaches the ***JALR***, from where it can deduce that ***opcode B*** at looplabel is the final retired instruction. However, there are circumstances where this approach does not work. For example, consider the case where there is an exception at looplabel + 4. In this case, the decoder cannot tell whether this occurred during the 1st or 2nd loop iterations, without additional information from the encoder. This is the purpose of the **updiscon** field. In more detail:

There are four scenarios to consider:

1. Code executes through to the end of the 1st loop iteration, and the encoder reports looplabel using format 1/2 following the ***JALR***, then carries on executing the 2nd pass of the loop. In this case **updiscon == notify**. The next packet will be a format 1/2;

2. Code executes through to the end of the 1st loop iteration, but there is an exception, privilege change or resync at the instruction following the ***JALR*** (i.e. at looplabel + 4). In this case, the encoder reports looplabel using format 1/2 following the ***JALR***, with **updiscon == !notify**, and the next packet is a format 3;

3. An exception occurs after the 1st execution of looplabel. In this case, the encoder reports looplabel using format 0/1/2 and again, **updiscon == notify**, and the next packet is a format 3;

4. The hart requests the encoder to notify retirement of the instruction at looplabel. In this case, the encoder reports the 1st execution of looplabel with **notify == !address[MSB]**, and subsequent executions with **notify == address[MSB]**.

Looking at this from the perspective of the decoder, the decoder receives a format 1/2 reporting the address of the 1st instruction in the loop (looplabel). It follows the execution path from the last reported address, until it reaches looplabel. Because looplabel is not preceded by an uninferable discontinuity, it must take the value of **notify** and **updiscon** into consideration, and may need to wait for the next packet in order to determine whether it has reached the final retired instruction:

- If **updiscon == !notify**, this indicates case 2. The decoder must continue until it encounters looplabel a 2nd time;

- If **updiscon == notify**, the decoder cannot yet distinguish cases 1 and 3, and must wait for the next packet.

  - If the next packet is a format 3, this is case 3. The decoder has already reached the correct instruction;

  - If the next packet is a format 1/2, this is case 1. The decoder must continue until it encounters looplabel a 2nd time.

- If **notify == !address[MSB]**, this indicates case 4, 1st iteration. The decoder has reached the correct instruction.

This example uses an exception at looplabel + 4, but anything that could cause a format 3 for looplabel + 4 would result in the same behavior: a privilege change, or the expiry of the resync timer. It could also occur if looplabel was the last traced instruction (because tracing was disabled for some reason). See section 6.5.1 for further discussion of this point.

Note: Correct decoder behavior could have been achieved by implementing the **notify** bit only, setting it to the inverse of **address[MSB]** whenever an address is reported and it is not the instruction following an uninferable discontinuity. However, this would have been much less efficient, as this would have required **notify** to be different from **address[MSB]** the majority of the time when outputting a format 1/2 before an exception, interrupt or resync (as the probability of this instruction being the target of an uninferable jump is low). Using 2 separate bits results in superior compression.

DISCUSSION POINT: The **notify** bit will only be different from **address[MSB]** when the hart has requested notification. As such, it will have no impact on 'normal' encoding efficiency. However, if the instruction being notified occurs in a loop, it may result in a relatively high number of notification packets close together in time. Because the **address** is typically 32-bits or more in length, but differtial address values are typically small, a notification packet is significantly less efficient. This could be mitigated by partitioning **address** into two parts, with **notify** in between (for example, **address[15:0]**, **notify**, **address[MSB:16]**). This would signicantly reduce

the efficiency impact of sending a notification. In return, the efficiency of a 'normal' format 1/2 packet may be slightly reduced if the differential address value being reported requires more than 16 bits, as some of the address bits are now after the **notify** bit. However, if the likelihood of such long jumps is low, this is probably worth it.

6.6.3   Format 2 irfail and irdepth fields

These bits are encoded so that most of the time they will take the same value as the **updiscon** field, and will therefore compress away, having no impact on the encoding efficiency. If implicit_return mode is enabled, and the encoder maintains a stack of predicted return addresses that are compared with the actual return addresses, then a *te_inst* packet will be generated if a misprediction occurs. In order to correctly reconstruct the execution path of the program, the decoder will need to know which return it was that failed. If a return is reported because the return address stack is empty, these fields will take the same value as the **updiscon** field.

6.7   Format 1 packets

This packet branch information, and is used when either the branch information must be reported (for example because the branch map is full), or whe the address of an instruction must be reported, and there has been at least one branch since the previous packet. If included, the address is in differential format unless full address mode is enabled (see section 2.2.2).

If branch prediction is supported and is enabled, then there is a choice of whether to output a full branch map, or a count of correctly predicted branches. The count format is used if the number of correctly predicted branches is at least 31. If there are 31 unreported branches (i.e. the branch map is full), but not all of them were predicted correctly, then the branch map will be output. A branch count will be output under the following conditions:

- A branch is mis-predicted. The count value will be the number of correctly predicted branches, minus 31. No address information is provided - it is implicitly that of the branch which failed prediction;

- An updiscon, interrupt or exception requires the encoder to output an address. In this case the encoder will output the branch count (number of correctly predicted branches, minus 31);

- The branch count reaches its maximum value. Strictly speaking an address isn't required for this case, but is included to avoid having to distinguish the packet format from the case above. t will occur so rarely that the bandwidth impact can be ignored.

6.7.1   Format 1 updiscon field

See section 6.6.2.

Table 6.6: Packet format 1 - address, branch map

Field name	Bits	Description
format	2	01 (diff-delta): includes branch information and may include differential address
branches	5	Number of valid bits **branch_map**. The length of branch-map is determined as follows: 0: (cannot occur for this format) 1: 1 bit 2-3: 3 bits 4-7: 7 bits 8-15: 15 bits 16-31: 31 bits For example if branches = 12, **branch_map** is 15 bits long, and the 12 LSBs are valid.
branch_map	Determined by **branches** field.	An array of bits indicating whether branches are taken or not. Bit 0 represents the oldest branch instruction executed. For each bit: 0: branch taken 1: branch not taken
address	$iaddress\_width\_p$ - $iaddress\_lsb\_p$	Differential instruction address.
notify	1	If the value of this bit is different from the MSB of **address**, it indicates that this packet is reporting an instruction that is not the target of an uninferable discontinuity because a notification was requested via **context_type** (see table 3.1).
updiscon	1	If the value of this bit is different from the MSB of **notify**, it indicates that this packet is reporting the instruction following an uninferable discontinuity and is also the instruction before an exception, privilege change or resync (i.e. it will be followed immediately by a format 3 *te_inst*).
irfail	1	If the value of this bit is different from **updiscon**, it indicates that this packet is reporting the instruction following a return because its address differs from the predicted return address at the top of the implicit_return return address stack.
irdepth	$return\_stack\_size\_p$	If the value of **irfail** is different from **updiscon**, this field indicates the number of entries on the return address stack (i.e. the entry number of the return that failed). If **irfail** is the same value as **updiscon**, all bits in this field will also be the same value as **updiscon**.

Table 6.7: Packet format 1 - no address, branch map

Field name	Bits	Description
format	2	01 (diff-delta): includes branch information and may include differential address
branches	5	Number of valid bits in **branch_map**. The length of **branch_map** is determined as follows: 0: 31 bits, no **address** in packet 1-31: (cannot occur for this format)
branch_map	31	An array of bits indicating whether branches are taken or not. Bit 0 represents the oldest branch instruction executed. For each bit: 0: branch taken 1: branch not taken
branch_fmt	2	Both bits set to the same value as **branch_map[30]** indicates that the preceding field is **branch_map**.

Table 6.8: Packet format 1 - no address, branch count

Field name	Bits	Description
format	2	01 (diff-delta): includes branch information and may include differential address
branches	5	Number of valid bits in **branch_map**. The length of **branch_map** is determined as follows: 0: 31 bits, no **address** in packet 31-1: (cannot occur for this format)
branch_count	31	Count of the number of correctly predicted branches, minus 31.
branch_fmt	2	Set to 01, indicates that the packet contains a **branch_count** field, no **address** field, and that the next branch failed prediction.

Table 6.9: Packet format 1 - address, branch count

Field name	Bits	Description
format	2	01 (diff-delta): includes branch information and may include differential address
branches	5	Number of valid bits in **branch_map**. The length of **branch_map** is determined as follows: 0: 31 bits, no **address** in packet 31-1: (cannot occur for this format)
branch_count	31	Count of the number of correctly predicted branches, minus 31.
branch_fmt	2	Set to 10, indicates that the packet contains a **branch_count** field and an **address** field. This will be the case if the packet is output because it is necessary to report an address (e.g. following an updiscon, or if the next instruction is an exception), or because **branch_count** has reached 0xffff).
bpsuccess	1	Set to 1 if the address points to a branch instruction and the branch was predicted correctly.
address	*iaddress_width_p - iaddress_lsb_p*	Differential instruction address.
notify	1	If the value of this bit is different from the MSB of **address**, it indicates that this packet is reporting an instruction that is not the target of an uninferable discontinuity because a notification was requested via **context_type** (see table 3.1).
updiscon	1	If the value of this bit is different from **notify**, it indicates that this packet is reporting the instruction following an uninferable discontinuity and is also the instruction before an exception, privilege change or resync (i.e. it will be followed immediately by a format 3 *te_inst*).
irfail	1	If the value of this bit is different from **updiscon**, it indicates that this packet is reporting the instruction following a return because its address differs from the predicted return address at the top of the implicit_return return address stack.
irdepth	*return_stack_size_p*	If the value of **irfail** is different from **updiscon**, this field indicates the number of entries on the return address stack (i.e. the entry number of the return that failed). If **irfail** is the same value as **updiscon**, all bits in this field will also be the same value as **updiscon**.

6.7.2   Format 1 branch_map field

When the branch map becomes full it must be reported, but in most cases there is no need to report an address. This is indicated by setting **branches** to 0. The exception to this is when the instruction immediately prior to the final branch causes an uninferable discontinuity, in which case **branches** is set to 31.

The choice of sizes (1, 3, 7, 15, 31) is designed to minimize efficiency loss. On average there will be some 'wasted' bits because the number of branches to report is less than the selected size of the **branch_map** field. Using a tapered set of sizes means that the number of wasted bits will on average be less for shorter packets. If the number of branches between updiscons is randomly distributed then the probability of generating packets with large branch counts will be lower, in which case increased waste for longer packets will have less overall impact. Furthermore, the rate at which packets are generated can be higher for lower branch counts, and so reducing waste for this case will improve overall bandwidth at times where it is most important.

6.7.3   Format 1 branch_fmt field

This is encoded so that when reporting a branch map it will take the same value as the MSB of the **branch_map** field, so that extra bits are only required when reporting predicted branch counts, and reporting a branch map is unaffected. Even for the most pathological case (32 correctly predicted branches followed by a misprediction), the total number of bits used is still fewer than if using just the branch map format.

6.7.4   Format 1 bpsuccess field

When a branch count is reported without an address it is because a branch has failed the prediction. However, when an address is reported along with a branch count, it will be because the packet was initiated by an uninferable discontinuity, an exception, or because a branch has been encountered when the number of correctly predicted branches is 0xffff. For the latter case, the reported address will always be for a branch, and in the former cases it may be. If it is a branch, it is necessary to be explicit about whether or not the prediction was met or not. If it is met, then the reported address is that of the last correctly predicted branch.

6.7.5   Format 1 irfail and irdepth fields

See section 6.7.5.

6.8   Format 0 packets

This format is intended for optional efficiency extensions. Currently only one extension is defined, for reporting the jump target cache index.

Table 6.10: Packet format 0, subformat 0 - jump target index, branch map

Field name	Bits	Description
format	2	00 (opt-ext): formats for optional efficiency extensions
subformat	See section 6.8.1	0 (jump target cache)
index	*cache_size_p*	Jump target cache index of entry containing target address.
branches	5	Number of valid bits in **branch_map**. The length of **branch_map** is determined as follows: 0: (cannot occur for this format) 1: 1 bit 2-3: 3 bits 4-7: 7 bits 8-15: 15 bits 16-31: 31 bits For example if branches = 12, **branch_map** is 15 bits long, and the 12 LSBs are valid.
branch_map	Determined by **branches** field.	An array of bits indicating whether branches are taken or not. Bit 0 represents the oldest branch instruction executed. For each bit: 0: branch taken 1: branch not taken
irfail	1	If the value of this bit is different from **branch_map[MSB]**, it indicates that this packet is reporting the instruction following a return because its address differs from the predicted return address at the top of the implicit_return return address stack.
irdepth	*return_stack_size_p*	If the value of **irfail** is different from **branch_map[MSB]**, this field indicates the number of entries on the return address stack (i.e. the entry number of the return that failed). If **irfail** is the same value as **branch_map[MSB]**, all bits in this field will also be the same value as **branch_map[MSB]**.

Table 6.11: Packet fFormat 0, subformat 0 - jump target index, no branch map

Field name	Bits	Description
format	2	00 (opt-ext): formats for optional efficiency extensions
subformat	See section 6.8.1	0 (jump target cache)
index	*cache_size_p*	Jump target cache index of entry containing target address.
branches	5	Number of valid bits in **branch_map**. The length of **branch_map** is determined as follows: 0: no **branch_map** in packet 1-31: (cannot occur for this format)
irfail	1	If the value of this bit is different from **branches[MSB]**, it indicates that this packet is reporting the instruction following a return because its address differs from the predicted return address at the top of the implicit_return return address stack.
irdepth	*return_stack_size_p*	If the value of **irfail** is different from **branches[MSB]**, this field indicates the number of entries on the return address stack (i.e. the entry number of the return that failed). If **irfail** is the same value as **branches[MSB]**, all bits in this field will also be the same value as **branches[MSB]**.

6.8.1   Format 0 subformat field

The width of this field depends on the number of optional formats supported. Currently, only one optional format is defined (branch target cache), and as such the width of this field is 0 (i.e. it is omitted). The width is specified by the *f0s_width* discovery field (see section 7.1). Provision of this field allows additional formats to be added in future without reducing the efficiency of the existing formats.

6.8.2   Format 0 irfail and irdepth fields

These bits are encoded so that most of the time they will take the same value as the **branch_map[MSB]** or **branches[MSB]** field. Purpose and behaviour is as described in section 6.7.5. They are included to allow return addresses that fail the implicit return prediction but which reside in the jump target cache to be reported using this format. An implementation could omit these if all implicit return failures are reported using format 1.

Chapter 7

Parameters and Discovery

This document defines a number of parameters for describing aspects of the encoder such as the widths of buses, the presence or absence of optional features and the size of resources, as listed in Table 7.1.

Depending on the implementation, some parameters may be inherently fixed whilst others may be passed in to the design by some means.

Table 7.1: Parameters to the encoder

Parameter name	Range	Description
arch_p		The architecture specification version with which the encoder is compliant (0 for initial version).
bpred_size_p		Number of entries in the branch predictor is $2^{\textbf{bpred\_size\_p}}$. Minimum number of entries is 2, so a value of 0 indicates that there is no branch predictor implemented.
cache_size_p		Number of entries in the jump target cache is $2^{\textbf{cache\_size\_p}}$. Minimum number of entries is 2, so a value of 0 indicates that there is no jump target cache implemented.
call_counter_size_p		Number of bits in the nested call counter is $2^{\textbf{call\_counter\_size\_p}}$. Minimum number of entries is 2, so a value of 0 indicates that there is no implicit return call counter implemented.
context_type_width_p		Width of the **context_type** bus
context_width_p		Width of context bus
ecause_width_p		Width of exception cause bus
ecause_choice_p		Number of bits of exception cause to match using multiple choice
f0s_width_p		Width of the **subformat** field in format 0 te_inst packets (see section 6.8.1).
filter_context_p	0 or 1	Filtering on context supported when 1
filter_ecause_p		Filtering on exception cause supported when non_zero. Number of nested exceptions supported is $2^{\text{filter\_ecause\_p}}$
filter_interrupt_p	0 or 1	Filtering on interrupt supported when 1
filter_privilege_p	0 or 1	Filtering on privilege supported when 1
filter_tval_p	0 or 1	Filtering on trap value supported when 1
iaddress_lsb_p		LSB of instruction address bus to trace. 1 is compressed instructions are supported, 2 otherwise
iaddress_width_p		Width of instruction address bus. This is the same as DXLEN
iretire_width_p		Width of the **iretire** bus
ilastsize_width_p		Width of the **ilastsize** bus
itype_width_p		Width of the **itype** bus
nocontext_p	0 or 1	Exclude context from te_inst packets if 1
ntkn_width_p		Width of the **ntkn** bus
privilege_width_p		Width of privilege bus
retires_p		Maximum number of instructions that can be retired per block
return_stack_size_p		Number of entries in the return address stack is $2^{\textbf{return\_stack\_size\_p}}$. Minimum number of entries is 2, so a value of 0 indicates that there is no implicit return stack implemented.
sijump_p	0 or 1	**sijump** is used to identify sequentially inferable jumps
taken_branches_p		Number of times **iretire**, **itype**, **ntkn** is replicated
impdef_width_p		Width of implementation-defined input bus

7.1 Discovery of encoder parameters

To operate correctly, the decoder must be able to determine some of the encoder's parameters at runtime, in the form of discoverable attributes. These parameters must be discoverable by the decoder, or else be fixed at the default value (in other words, if an encoder does not make a particular parameter discoverable, it must implement only the default value of that parameter, which the decoder will also use). Table 7.2 lists the required discoverable attributes.

To access the discoverable attributes, some external entity, for example a debugger or a supervisory hart, must request it from the encoder. The encoder will provide the discovery information in one or more different formats. The preferred format is a packet which is sent over the trace infrastructure. Another format would be allowing the external entity to read the values from some register or memory mapped space maintained by the encoder. Section 7.2 gives an example of how this may be accomplished.

Table 7.2: Required attributes

Name	Default	Parameter mapping
arch	0	*arch_p*
bpred_size	0	*bpred_size_p*
cache_size	0	*cache_size_p*
call_counter_size	0	*call_counter_size_p*
context_width	0	*context_width_p* - 1
ecause_width	3	*ecause_width_p* - 1
f0s_width	0	*f0s_width_p*
iaddress_lsb	0	*iaddress_lsb_p* - 1
iaddress_width	31	*iaddress_width_p* - 1
nocontext	0	*nocontext*
privilege_width	2	*privilege_width_p* - 1
return_stack_size	0	*return_stack_size_p*
sijump	0	*sijump_p*

For ease of use it is further recommended that all of the encoder's parameters be mapped to discoverable attributes, even if not directly required by the decoder. In particular, attributes related to filtering capabilities. Table 7.3 lists the attributes associated with the filtering recommendations discussed in Chapter 4, and Table 7.4 lists attributes related to other parameters mentioned in this document.

Table 7.3: Optional filtering attributes

Name	Default	Parameter mapping
comparators	0	*comparators_p* - 1
filters	0	*filters_p* - 1
ecause_choice	5	*ecause_choice_p*
filter_context	1	*filter_context_p*
filter_ecause	1	*filter_ecause_p*
filter_interrupt	1	*filter_interrupt_p*
filter_privilege	1	*filter_privilegep*
filter_tval	1	*filter_tval_p*

Table 7.4: Other recommended attributes

Name	Default	Description
context_type_width	1	*context_type_width_p* - 1
ilastsize_width	0	*ilastsize_width_p* - 1
itype_width	4	*itype_width_p* - 1
iretire_width	2	*iretire_width_p* - 1
ntkn_width	0	*ntkn_width_p* - 1
retires	0	*retires_p* - 1
taken_branches	0	*taken_branches_p* - 1
impdef_width	0	*impdef_width_p* - 1

7.2 Example ipxact description

This section provides an example of discovery information represented in the ipxact form.

```xml
<?xml version="1.0" encoding="UTF-8"?>
<ipxact:component
   xmlns:ipxact="http://www.accellera.org/XMLSchema/IPXACT/1685-2014"
   xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
   xsi:schemaLocation="http://www.accellera.org/XMLSchema/IPXACT/1685-2014
                       http://www.accellera.org/XMLSchema/IPXACT/1685-2014/index.xsd">
   <ipxact:vendor>UltraSoC</ipxact:vendor>
   <ipxact:library>TraceEncoder</ipxact:library>
   <ipxact:name>TraceEncoder</ipxact:name>
   <ipxact:version>0.8</ipxact:version>
   <ipxact:memoryMaps>
      <ipxact:memoryMap>
         <ipxact:name>Trace Encoder Register Map</ipxact:name>
         <ipxact:addressBlock>
            <ipxact:name>>Trace Encoder Register Address Block</ipxact:name>
            <ipxact:baseAddress>0</ipxact:baseAddress>
            <ipxact:range>128</ipxact:range>
            <ipxact:width>64</ipxact:width>
```

```
<ipxact:register>
   <ipxact:name>discovery_info_0</ipxact:name>
   <ipxact:addressOffset>'h0</ipxact:addressOffset>
   <ipxact:size>64</ipxact:size>
   <ipxact:access>read-only</ipxact:access>
   <ipxact:field>
      <ipxact:name>version</ipxact:name>
      <ipxact:description>text</ipxact:description>
      <ipxact:bitOffset>0</ipxact:bitOffset>
      <ipxact:bitWidth>4</ipxact:bitWidth>
   </ipxact:field>
   <ipxact:field>
      <ipxact:name>minor_revision</ipxact:name>
      <ipxact:description>text</ipxact:description>
      <ipxact:bitOffset>4</ipxact:bitOffset>
      <ipxact:bitWidth>4</ipxact:bitWidth>
   </ipxact:field>
   <ipxact:field>
      <ipxact:name>arch</ipxact:name>
      <ipxact:description>text</ipxact:description>
      <ipxact:bitOffset>8</ipxact:bitOffset>
      <ipxact:bitWidth>4</ipxact:bitWidth>
   </ipxact:field>
   <ipxact:field>
      <ipxact:name>bpred_size</ipxact:name>
      <ipxact:description>text</ipxact:description>
      <ipxact:bitOffset>12</ipxact:bitOffset>
      <ipxact:bitWidth>4</ipxact:bitWidth>
   </ipxact:field>
   <ipxact:field>
      <ipxact:name>cache_size</ipxact:name>
      <ipxact:description>text</ipxact:description>
      <ipxact:bitOffset>16</ipxact:bitOffset>
      <ipxact:bitWidth>4</ipxact:bitWidth>
   </ipxact:field>
   <ipxact:field>
      <ipxact:name>call_counter_size</ipxact:name>
      <ipxact:description>text</ipxact:description>
      <ipxact:bitOffset>20</ipxact:bitOffset>
      <ipxact:bitWidth>3</ipxact:bitWidth>
   </ipxact:field>
   <ipxact:field>
      <ipxact:name>comparators</ipxact:name>
      <ipxact:description>text</ipxact:description>
      <ipxact:bitOffset>23</ipxact:bitOffset>
      <ipxact:bitWidth>3</ipxact:bitWidth>
   </ipxact:field>
   <ipxact:field>
```

```
        <ipxact:name>context_type_width</ipxact:name>
        <ipxact:description>text</ipxact:description>
        <ipxact:bitOffset>26</ipxact:bitOffset>
        <ipxact:bitWidth>5</ipxact:bitWidth>
    </ipxact:field>
    <ipxact:field>
        <ipxact:name>context_width</ipxact:name>
        <ipxact:description>text</ipxact:description>
        <ipxact:bitOffset>31</ipxact:bitOffset>
        <ipxact:bitWidth>5</ipxact:bitWidth>
    </ipxact:field>
    <ipxact:field>
        <ipxact:name>ecause_choice</ipxact:name>
        <ipxact:description>text</ipxact:description>
        <ipxact:bitOffset>36</ipxact:bitOffset>
        <ipxact:bitWidth>3</ipxact:bitWidth>
    </ipxact:field>
    <ipxact:field>
        <ipxact:name>ecause_width</ipxact:name>
        <ipxact:description>text</ipxact:description>
        <ipxact:bitOffset>39</ipxact:bitOffset>
        <ipxact:bitWidth>4</ipxact:bitWidth>
    </ipxact:field>
    <ipxact:field>
        <ipxact:name>filters</ipxact:name>
        <ipxact:description>text</ipxact:description>
        <ipxact:bitOffset>43</ipxact:bitOffset>
        <ipxact:bitWidth>4</ipxact:bitWidth>
    </ipxact:field>
    <ipxact:field>
        <ipxact:name>filter_context</ipxact:name>
        <ipxact:description>text</ipxact:description>
        <ipxact:bitOffset>47</ipxact:bitOffset>
        <ipxact:bitWidth>1</ipxact:bitWidth>
    </ipxact:field>
    <ipxact:field>
        <ipxact:name>filter_ecause</ipxact:name>
        <ipxact:description>text</ipxact:description>
        <ipxact:bitOffset>48</ipxact:bitOffset>
        <ipxact:bitWidth>4</ipxact:bitWidth>
    </ipxact:field>
    <ipxact:field>
        <ipxact:name>filter_interrupt</ipxact:name>
        <ipxact:description>text</ipxact:description>
        <ipxact:bitOffset>52</ipxact:bitOffset>
        <ipxact:bitWidth>1</ipxact:bitWidth>
    </ipxact:field>
    <ipxact:field>
```

```
            <ipxact:name>filter_privilege</ipxact:name>
            <ipxact:description>text</ipxact:description>
            <ipxact:bitOffset>53</ipxact:bitOffset>
            <ipxact:bitWidth>1</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
            <ipxact:name>filter_tval</ipxact:name>
            <ipxact:description>text</ipxact:description>
            <ipxact:bitOffset>54</ipxact:bitOffset>
            <ipxact:bitWidth>1</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
            <ipxact:name>filter_impdef</ipxact:name>
            <ipxact:description>text</ipxact:description>
            <ipxact:bitOffset>55</ipxact:bitOffset>
            <ipxact:bitWidth>1</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
            <ipxact:name>f0s_width</ipxact:name>
            <ipxact:description>text</ipxact:description>
            <ipxact:bitOffset>56</ipxact:bitOffset>
            <ipxact:bitWidth>2</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
            <ipxact:name>iaddress_lsb</ipxact:name>
            <ipxact:description>text</ipxact:description>
            <ipxact:bitOffset>58</ipxact:bitOffset>
            <ipxact:bitWidth>2</ipxact:bitWidth>
        </ipxact:field>
    </ipxact:register>


    <ipxact:register>
        <ipxact:name>discovery_info_1</ipxact:name>
        <ipxact:addressOffset>'h4</ipxact:addressOffset>
        <ipxact:size>64</ipxact:size>
        <ipxact:access>read-only</ipxact:access>
        <ipxact:field>
            <ipxact:name>iaddress_width</ipxact:name>
            <ipxact:description>text</ipxact:description>
            <ipxact:bitOffset>0</ipxact:bitOffset>
            <ipxact:bitWidth>7</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
            <ipxact:name>ilastsize_width</ipxact:name>
            <ipxact:description>text</ipxact:description>
            <ipxact:bitOffset>7</ipxact:bitOffset>
            <ipxact:bitWidth>7</ipxact:bitWidth>
```

```
        </ipxact:field>
        <ipxact:field>
           <ipxact:name>itype_width</ipxact:name>
           <ipxact:description>text</ipxact:description>
           <ipxact:bitOffset>14</ipxact:bitOffset>
           <ipxact:bitWidth>7</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
           <ipxact:name>iretire_width</ipxact:name>
           <ipxact:description>text</ipxact:description>
           <ipxact:bitOffset>21</ipxact:bitOffset>
           <ipxact:bitWidth>7</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
           <ipxact:name>nocontext</ipxact:name>
           <ipxact:description>text</ipxact:description>
           <ipxact:bitOffset>28</ipxact:bitOffset>
           <ipxact:bitWidth>1</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
           <ipxact:name>ntkn_width</ipxact:name>
           <ipxact:description>text</ipxact:description>
           <ipxact:bitOffset>29</ipxact:bitOffset>
           <ipxact:bitWidth>2</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
           <ipxact:name>privilege_width</ipxact:name>
           <ipxact:description>text</ipxact:description>
           <ipxact:bitOffset>31</ipxact:bitOffset>
           <ipxact:bitWidth>2</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
           <ipxact:name>retires</ipxact:name>
           <ipxact:description>text</ipxact:description>
           <ipxact:bitOffset>33</ipxact:bitOffset>
           <ipxact:bitWidth>3</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
           <ipxact:name>return_stack_size</ipxact:name>
           <ipxact:description>text</ipxact:description>
           <ipxact:bitOffset>36</ipxact:bitOffset>
           <ipxact:bitWidth>4</ipxact:bitWidth>
        </ipxact:field>
        <ipxact:field>
           <ipxact:name>sijump</ipxact:name>
           <ipxact:description>text</ipxact:description>
           <ipxact:bitOffset>40</ipxact:bitOffset>
           <ipxact:bitWidth>1</ipxact:bitWidth>
```

```
                </ipxact:field>
                <ipxact:field>
                    <ipxact:name>taken_branches</ipxact:name>
                    <ipxact:description>text</ipxact:description>
                    <ipxact:bitOffset>41</ipxact:bitOffset>
                    <ipxact:bitWidth>4</ipxact:bitWidth>
                </ipxact:field>
                <ipxact:field>
                    <ipxact:name>impdef_width</ipxact:name>
                    <ipxact:description>text</ipxact:description>
                    <ipxact:bitOffset>45</ipxact:bitOffset>
                    <ipxact:bitWidth>5</ipxact:bitWidth>
                </ipxact:field>
            </ipxact:register>

        </ipxact:addressBlock>
        <ipxact:addressUnitBits>8</ipxact:addressUnitBits>
      </ipxact:memoryMap>
   </ipxact:memoryMaps>
</ipxact:component>
```

Chapter 8

Future Directions

The current focus is the compressed branch trace, however there a number of other types of processor trace that would be useful (detailed below in no particular order). These should be considered as possible features that maybe added in the future, once the current scope has been completed.

8.1 Data trace

The trace encoder will output packets to communicate information about loads and stores to an off-chip decoder. To reduce the amount of bandwidth required, reporting data values will be optional, and both address and data will be able to be encoded differentially when it is beneficial to do so. This entails outputting the difference between the new value and the previous value of the same transfer size, irrespective of transfer direction.

Unencoded values will be used for synchronisation and at other times.

8.2 Fast profiling

In this mode the encoder will provide a non-intrusive alternative to the traditional method of profiling, which requires the processor to be halted periodically so that the program counter can be sampled. The encoder will issue packets when an exception, call or return is detected, to report the next instruction executed (i.e. the destination instruction). Optionally, the encoder will also be able to report the current instruction (i.e. the source instruction).

8.3 Inter-instruction cycle counts

In this mode the encoder will trace where the hart is stalling by reporting the number of cycles between successive instruction retirements.

8.4   Transport

After the current charter has been satisfied the transport mechanism should be defined and standardised. This will include Aurora based serdes, PCIe and Ethernet.

Chapter 9

Decoder

This decoder implementation assumes there is no branch predictor or return address stack (*return_stack_size_p* and *bpred_size_p* both zero).

9.1   Decoder pseudo code

```
# global variables
global      pc                          # Reconstructed program counter
global      last_pc                     # PC of previous instruction
global      branches = 0                # Number of branches to process
global      branch_map = 0              # Bit vector of not taken/taken (1/0) status
                                        #   for branches
global bool  stop_at_last_branch = FALSE # Flag to indicate reconstruction is to end at
                                        #   the final branch
global bool  inferred_address = FALSE   # Flag to indicate that reported address from
                                        #   format 0/1/2 was not following an uninferrable
                                        #   jump (and is therefore inferred)
global bool  start_of_trace = TRUE      # Flag indicating 1st trace packet still
                                        #   to be processed
global      address                     # Reconstructed address from te_inst messages
global      options                     # Operating mode flags
global      call_counter = 0            # Count of number of nested calls being traced
global array return_stack               # Array holding return address stack
```

```
# Process te_inst packet.  Call each time a te_inst packet is received #
function process_te_inst (te_inst)
  if (te_inst.format == 3)
    inferred_address = FALSE
    address        = (te_inst.address << discovery_response.iaddress_lsb)
    if (te_inst.subformat == 3) # Support packet
      process_support(te_inst)
      return

    if (te_inst.subformat == 1 or start_of_trace)
      branches    = 0
      branch_map  = 0
    if (is_branch(get_instr(address))) # 1 unprocessed branch if this instruction is a branch
      branch_map = branch_map | (te_inst.branch << branches)
      branches++
    if (te_inst.subformat == 0 and !start_of_trace)
      follow_execution_path(address, te_inst)
    else
      pc            = address
      last_pc       = pc # previous pc not known but ensures correct
                        #  operation for is_sequential_jump()
    start_of_trace = FALSE
    call_counter   = 0

  else
    if (start_of_trace) # This should not be possible!
      ERROR: Expecting trace to start with format 3
      return
    if (te_inst.format == 2 or te_inst.branches != 0)
      stop_at_last_branch = FALSE
      if (options.full_address)
        address  = (te_inst.address << discovery_response.iaddress_lsb)
      else
        address += (te_inst.address << discovery_response.iaddress_lsb)
    if (te_inst.format == 1)
      stop_at_last_branch = (te_inst.branches == 0)
      # Branch map will contain <= 1 branch (1 if last reported instruction was a branch)
      branch_map = branch_map | (te_inst.branch_map << branches)
      if (te_inst.branches == 0)
        branches += 31
      else
        branches += te_inst.branches

    follow_execution_path(address, te_inst)
```

```
# Follow execution path to reported address #
function follow_execution_path(address, te_inst)

  local previous_address = pc
  while (TRUE)
    if (inferred_address) # iterate again from previously reported address to
                          #   find second occurrence
      next_pc(previous_address)
      if (pc == previous_address)
        inferred_address = FALSE
    else
      next_pc(address)
      if (branches == 1 and is_branch(get_instr(pc)) and stop_at_last_branch)
        # Reached final branch - stop here (do not follow to next instruction as
        #  we do not yet know whether it retires)
        stop_at_last_branch = FALSE
        return
      if (pc == address and is_uninferrable_discon(get_instr(last_pc)))
        # Reached reported address following an uninferrable discontinuity - stop here
        if (branches > (is_branch(get_instr(pc)) ? 1 : 0))
          # Check all branches processed (except 1 if this instruction is a branch)
          ERROR: unprocessed branches
        return
      if (te_inst.format != 3 and pc == address and
        (te_inst.notify != te_inst.address[MSB]) and
        (branches == (is_branch(get_instr(pc)) ? 1 : 0)))
          # All branches processed, and reached reported address due to notification,
          # not as an uninferrable jump target
        return
      if (te_inst.format != 3 and pc == address and
        (te_inst.updiscon == te_inst.notify) and
        (branches == (is_branch(get_instr(pc)) ? 1 : 0)))
          # All branches processed, and reached reported address, but not as an
          #   uninferrable jump target
          # Stop here for now, though flag indicates this may not be
          #  final retired instruction
        inferred_address = TRUE
        return
      if (te_inst.format == 3 and pc == address and
        (branches == (is_branch(get_instr(pc)) ? 1 : 0)))
        # All branches processed, and reached reported address
        return
```

```
# Compute next PC #
function next_pc (address)

  local instr   = get_instr(pc)
  local this_pc = pc

  if (is_inferrable_jump(instr))
    pc += instr.imm
  else if (is_sequential_jump(instr, last_pc)) # lui/auipc followed by
                                               #  jump using same register
    pc = sequential_jump_target(pc, last_pc)
  else if (is_implicit_return(instr))
    pc = pop_return_stack()
  else if (is_uninferrable_discon(instr))
    if (stop_at_last_branch)
      ERROR: unexpected uninferrable discontinuity
    else
      pc = address
  else if (is_taken_branch(instr))
    pc += instr.imm
  else
    pc += instruction_size(instr)

  if (is_call(instr))
    push_return_stack(this_pc)

  last_pc = this_pc

# Process support packet #
function process_support (te_inst)

  options = te_inst.options
    if (te_inst.qual_status != no_change)
      start_of_trace = TRUE # Trace ended, so get ready to start again
    if (te_inst.qual_status == ended_upd and inferred_address)
      local previous_address = pc
      inferred_address        = FALSE
      while (TRUE)
        next_pc(previous_address)
        if (pc == previous_address)
          return
    return
```

```
# Determine if instruction is a branch, adjust branch count/map,
#   and return taken status #
function is_taken_branch (instr)
  local bool taken = FALSE

  if (!is_branch(instr))
    return FALSE

  if (branches == 0)
    ERROR: cannot resolve branch
  else
    taken = !branch_map[0]
    branches--
    branch_map >> 1

  return taken

# Determine if instruction is a branch #
function is_branch (instr)

  if ((instr.opcode == BEQ)    or
      (instr.opcode == BNE)    or
      (instr.opcode == BLT)    or
      (instr.opcode == BGE)    or
      (instr.opcode == BLTU)   or
      (instr.opcode == BGEU)   or
      (instr.opcode == C.BEQZ) or
      (instr.opcode == C.BNEZ))
    return TRUE

  return FALSE

# Determine if instruction is an inferrable jump #
function is_inferrable_jump (instr)

  if ((instr.opcode == JAL)   or
      (instr.opcode == C.JAL) or
      (instr.opcode == C.J)    or
      (instr.opcode == JALR and instr.rs1 == 0))
    return TRUE

  return FALSE
```

```
# Determine if instruction is an uninferrable jump #
function is_uninferrable_jump (instr)

  if ((instr.opcode == JALR and instr.rs1 != 0) or
      (instr.opcode == C.JALR)                     or
      (instr.opcode == C.JR))
    return TRUE

  return FALSE

# Determine if instruction is an uninferrable discontinuity #
function is_uninferrable_discon (instr)

  if (is_uninferrable_jump(instr) or
      (instr.opcode == URET)      or
      (instr.opcode == SRET)      or
      (instr.opcode == MRET)      or
      (instr.opcode == DRET))
    return TRUE
    # Note: The exception reporting mechanism means it is not
    #   necessary to include
    # ECALL, EBREAK or C.EBREAK in this test

  return FALSE

# Determine if instruction is a sequentially inferrable jump #
function is_sequential_jump (instr, prev_addr)

  if (not (is_uninferrable_jump(instr) and options.sijump))
    return FALSE

  local prev_instr = get_instr(prev_addr)

  if((prev_instr.opcode == AUIPC) or
     (prev_instr.opcode == LUI)   or
     (prev_instr.opcode == C.LUI))
    return (instr.rs1 == prev_instr.rd)

  return FALSE
```

```
# Find the target of a sequentially inferrable jump #
function sequential_jump_target (addr, prev_addr)

  local instr      = get_instr(addr)
  local prev_instr = get_instr(prev_addr)
  local target     = 0

  if (prev_instr.opcode == AUIPC)
    target = prev_addr
  target += prev_instr.imm
  if (instr.opcode == JALR)
    target += instr.imm

  return target

# Determine if instruction is a call #
# - excludes tail calls as they do not push an address onto the return stack
function is_call (instr)

  if ((instr.opcode == JALR and instr.rd == 1) or
      (instr.opcode == C.JALR)                   or
      (instr.opcode == JAL  and instr.rd == 1) or
      (instr.opcode == C.JAL))
    return TRUE

  return FALSE

# Determine if instruction return address can be implicitly inferred #
function is_implicit_return (instr)

  if (options.implicit_return == 0) # Implicit return mode disabled
    return FALSE

  if ((instr.opcode == JALR and instr.rs1 == 1 and instr.rd == 0) or
      (instr.opcode == C.JR and instr.rs1 == 1))
    return (call_counter > 0)

  return FALSE
```

```
# Push address onto return stack #
function push_return_stack (address)

  local call_counter_max = 2**(discovery_response.call_counter_width + 2)
  local instr            = get_instr(address)
  local link             = address

  if (call_counter == call_counter_max)
    # Delete oldest entry from stack to make room for new entry added below
    call_counter--
    for (i = 0; i < call_counter; i++)
      return_stack[i] = return_stack[i+1]

  link += instruction_size(instr)

  return_stack[call_counter] = link
  call_counter++

  return

# Pop address from return stack #
function pop_return_stack ()

  local link = return_stack[call_counter]

  call_counter-- # function not called if call_counter is 0, so no need
                 #  to check for underflow
  return link
```

Chapter 10

Example code and packets

In the following examples **ret** is referred to as uninferable, this is only true if implicit-return mode is off

1. **Call to debug_printf(), from 80001a84, in main():**

   ```
   00000000800019e8 <main>:
        ........: ...
        80001a80: f6d42423              sw a3,-152(s0)
        80001a84: ef4ff0ef              jal x1,80001178 <debug_printf>
   ```

   PC: 80001a84 ->80001178
   The target of the **jal** is inferable, thus NO te_inst packet is sent.

   ```
   0000000080001178 <debug_printf>:
        80001178: 7139                  addi sp,sp,-64
        8000117a: ...
   ```

2. **Return from debug_printf():**

   ```
        80001186: ...
        80001188: 6121                  addi sp,sp,64
        8000118a: 8082                  ret
   ```

   PC: 8000118a ->80001a88
   The target of the **ret** is uninferable, thus a **te_inst** packet IS sent:
   te_inst[format=2 (ADDR_ONLY): address=0x80001a88, updiscon=0]

65

```
        80001a88: 00000597                auipc a1,0x0
        80001a8c: 65058593                addi a1,a1,1616 # 800020d8 <main+0x6f0>
```

3. **exiting from Func_2(), with a final taken branch, followed by a *ret***

```
00000000800010b6 <Func_2>:
    ........:   ....
    800010da: 4781                        li a5,0
    800010dc: 00a05863                    blez a0,800010ec <Func_2+0x36>
```

PC: 800010dc ->800010ec, add branch TAKEN to branch_map, but no packet sent yet.
branches = 0; branch_map = 0;
branch_map |= 0 «branches++;

```
        800010ec:   60e2                    ld      ra,24(sp)
        800010ee:   6442                    ld      s0,16(sp)
        800010f0:   64a2                    ld      s1,8(sp)
        800010f2:   853e                    mv      a0,a5
        800010f4:   6105                    addi    sp,sp,32
        800010f6:   8082                    ret
```

PC: 800010f6 ->80001b8a
The target of the *ret* is uninferrable, thus a ***te_inst*** packet is sent, with ONE branch in the
branch_map
te_inst[format=1 (DIFF_DELTA): branches=1, branch_map=0x0, address=0x80001b8a
(Δ=0xab0) updiscon=0]

```
00000000800019e8 <main>:
    ........:   ....
    80001b8a: f4442603                    lw      a2,-188(s0)
    80001b8e: ....
```

4. **3 branches, then a function return back to Proc_1()**

```
0000000080001100 <Proc_6>:
    ........:   ....
    80001112: c080                        sw s0,0(s1)
    80001114: 4785                        li a5,1
    80001116: 02f40463                    beq s0,a5,8000113e <Proc_6+0x3e>
```

PC: 80001116 ->8000111a, add branch NOT taken to branch_map, but no packet sent yet.
branches = 0; branch_map = 0; branch_map |= 1 «branches++;

```
8000111a: c81d                              beqz s0,80001150 <Proc_6+0x50>
```

PC: 8000111a ->8000111c, add branch NOT taken to branch_map, but no packet sent yet.
branch_map |= 1 «branches++;

```
8000111c:   4709                        li      a4,2
8000111e:   04e40063                    beq     s0,a4,8000115e <Proc_6+0x5e>
```

PC: 8000111e ->8000115e, add branch TAKEN to branch_map, but no packet sent yet.
branch_map |= 0 «branches++;

```
8000115e: 60e2                  ld ra,24(sp)
80001160: 6442                  ld s0,16(sp)
80001162: c09c                  sw a5,0(s1)
80001164: 64a2                  ld s1,8(sp)
80001166: 6105                  addi sp,sp,32
80001168: 8082                  ret
```

```
0000000800011d6 <Proc_1>:
    ........:   ....
    80001258: 00093783          ld a5,0(s2)
    8000125c: ....
```

PC: 80001168 ->80001258
The target of the **ret** is uninferrable, thus a **te_inst** packet is sent, with THREE branches
in the branch_map
te_inst[format=1 (DIFF_DELTA): branches=3, branch_map=0x3, address=0x80001258
(Δ=0x148), updiscon=0]

5. **A complex example with 2 branches, 2 jal, and a ret**

```
0000000800011d6 <Proc_1>:
    ........:   ....
    8000121c: 441c              lw a5,8(s0)
    8000121e: c795              beqz a5,8000124a <Proc_1+0x74>
```

PC: 8000121e ->8000124a, add branch TAKEN to branch_map, but no packet sent yet.
branches = 0; branch_map = 0;
branch_map |= 0 «branches++;

```
8000124a: 44c8              lw a0,12(s1)
8000124c: 4799              li a5,6
8000124e: 00c40593          addi a1,s0,12
80001252: c81c              sw a5,16(s0)
80001254: eadff0ef          jal x1,80001100 <Proc_6>
```

PC: 80001254 ->80001100
The target of the ***jal*** is inferrable, thus no ***te_inst*** packet needs be sent.

```
0000000080001100 <Proc_6>:
80001100: 1101                      addi sp,sp,-32
80001102: e822                      sd s0,16(sp)
80001104: e426                      sd s1,8(sp)
80001106: ec06                      sd ra,24(sp)
80001108: 842a                      mv s0,a0
8000110a: 84ae                      mv s1,a1
8000110c: fedff0ef                  jal x1,800010f8 <Func_3>
```

PC: 8000110c ->800010f8
The target of the ***jal*** is inferrable, thus no ***te_inst*** packet needs to be sent.

```
00000000800010f8 <Func_3>:
800010f8: 1579                      addi a0,a0,-2
800010fa: 00153513                  seqz a0,a0
800010fe: 8082                      ret
```

PC: 800010fe ->80001110
The target of the ***ret*** is uninferrable, thus a ***te_inst*** packet will be sent shortly.

```
0000000080001100 <Proc_6>:
 ........:   ....
80001110: c115                      beqz a0,80001134 <Proc_6+0x34>
80001112: ....
```

PC: 80001110 ->80001112, add branch NOT TAKEN to branch_map.
branch_map |= 1 «branches++;
te_inst[format=1 (DIFF_DELTA): branches=2, branch_map=0x2, address=0x80001110
(Δ=0xfffffffffffffef4), updiscon=1]